

BLOOM-WILT: Black-Box Elicitation for Automated LLM Evaluation via Logit Steering

Anonymous Authors¹

Abstract

Stub.

1. Introduction

As language models are deployed into increasingly agentic settings, what matters for safety is less their average behaviour on a static benchmark than their worst-case behaviour on the long tail of *natural* inputs — the realistic user turns a model actually meets, not engineered adversarial strings (Jones et al., 2025). This is a question of *propensity* rather than adversarial robustness: a jailbreak forces a harmful string by pushing the target out of distribution, whereas we ask how often an ordinary conversation *happens* to elicit an undesired behaviour. A model that shrugs off every gibberish-suffix attack can still be drawn into racial bias, strategic deception, or self-preservation by dialogue no one would flag as adversarial — and it is these failures that surface once a model is deployed. Automated auditors such as BLOOM (Gupta et al., 2025) make finding them scalable: from a one-sentence description of a behaviour, an auditor model invents scenarios, converses with the target over several turns, and a judge scores the transcripts — reproducibly and with no human in the loop. And auditing a frontier model is necessarily *black-box*: its weights are unreleased, so an auditor reads the target’s outputs but never its activations, gradients, or parameters, which rules out white-box steering, ablation, or fine-tuning the target. Far from a handicap, this is the stronger setting in practice — benchmarking auditing agents on models with implanted behaviours, Sheshadri et al. (2026) find they perform *best* with black-box tools.

The bottleneck is how hard the auditor can push the target toward the behaviour. BLOOM only *samples* the target, but undesired behaviours live in the tail of its output distribution, so naive decoding needs many rollouts to surface

one clear example and underestimates how often the behaviour occurs (Beyer et al., 2026; Wu & Hilton, 2024). The obvious way to push harder is to borrow from the *jail-breaking* literature, but each of its tools forfeits something an auditor depends on. Token-level input attacks (Zou et al., 2024; Sadasivan et al., 2024) yield gibberish or off-scenario prompts no real user would type. Decoding-time logit arithmetic steers with a *second* model: a trained proxy that may be too weak to display a sophisticated behaviour itself (Zhao et al., 2024; Liu et al., 2024), or a pre-alignment checkpoint that has to be available (Aranguri & McGrath, 2025; Zhou et al., 2024); and learned investigators (Chowdhury et al., 2025) spend a reinforcement-learning run per behaviour. All of them, moreover, chase *any* output that elicits the behaviour, not one that stays probable under the unmodified target — the very thing a propensity estimate needs.

We close this gap with **WILT** (*Within-model Importance-weighted Logit Tilting*), which gives the auditor its missing leverage on the *output* side: rather than optimising the input, it steers the target’s own distribution toward the behaviour as it decodes. The target drafts a reply; we then resample each token from a logit-space mixture of its true distribution and a *corruption* distribution — the *same model*, now prompted to rewrite that draft so the behaviour appears — with a refusal direction optionally subtracted so it cannot simply fall back on refusing. Because the corruptor is the target itself, WILT needs no second model, no training and no checkpoint; the steering is exactly as capable as the audited model; and it reads nothing but output token distributions, so it stays fully black-box. A compliance prefill and that subtracted refusal term drive even safety-tuned targets to corrupt themselves, while a naturalness floor, repetition-aware selection, and a bank of diverse corruption prompts keep the output fluent and on-policy. The result is *guided importance sampling*: from the same one-sentence behaviour description, the audit surfaces behaviours that naive decoding would rarely reach in far fewer samples, yet each elicited transcript stays high-probability under the unmodified target and is scored under that same distribution — a faithful “this could really happen,” not capability inflation.

Contributions.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.




Figure 1. We search a tree of conversational trajectories, branching on both sides of the conversation and scoring each fan-out against the target’s own probabilities. Input-side branches are scored by target log-prefix-probability; output-side branches by a contrastive jailbroken-vs-safe logit gap. A judge scores each completed trajectory and a refiner updates the evaluator strategy for the next round.

- **BLOOM-WILT.** An extension of the BLOOM auditing pipeline that supplies elicitation pressure in the black-box setting: it changes only the decoder inside BLOOM’s rollout — leaving ideation, scenario generation, and judgement untouched — and steers the target toward any named behaviour in logit space, with a *self-corruption* distribution (the same model rewriting its own draft). It is training-free, uses one model and no checkpoint, and reads only output distributions.
- **Refusal-resistant by construction.** A subtracted refusal direction and a compliance prefix drive even safety-tuned targets to corrupt themselves, using only the target itself — no ablated, jailbroken, or fine-tuned proxy.
- **Natural, on-policy outputs.** A naturalness floor, repetition-aware selection, and a bank of diverse rewrite prompts keep every elicited transcript fluent and high-probability under the *unmodified* target.
- **Sample-efficient.** As guided importance sampling, WILT surfaces rare behaviours in far fewer samples than naive decoding and yields concrete transcripts — usable for inspection or adversarial training — rather than only a probability estimate.
- **Empirical gains across models and behaviours.**

Across four open-weight targets (Llama-3.2-3B, Phi-4-mini, Qwen3.6-4B, and Gemma-4-E4B) and six behaviours at matched compute, it lifts mean behaviour-presence from 2.4 to 6.5 on a 0–10 scale over the strongest baseline while holding the mean per-token target probability near 0.8, beating naive best-of- N sampling (BoN), weak-to-strong jailbreaking (W2S), and the input-side methods BEAST, FLRT, and a generalised PAIR (G-PAIR) on *both* elicitation and output probability. Per-component ablations isolate each ingredient above, and on a standard harmful-instruction benchmark WILT matches the attack-success protocol of the jailbreak literature.

2. Related Work

Automated behavioural evaluation. Perez et al. (2022) introduced the LLM-evaluator-against-target paradigm, automating red-teaming by having one language model generate test inputs for another and a classifier score the resulting transcripts. BLOOM (Gupta et al., 2025) formalises this into a four-stage pipeline (understanding, ideation, rollout, judgement) producing reproducible evaluation suites from a behaviour description. Petri (Fronsdal et al., 2025), the most widely used auditing agent, instead takes an *open-ended* approach — exploring for any concerning behaviour,

with tools to set system prompts, roll back conversations, and prefill responses — rather than targeting a single named behaviour as we do. Huang et al. (2025) benchmark static, offline, and online methods for multi-turn behaviour elicitation, finding that static test cases saturate within a year while online methods keep finding failures. We build on BLOOM: its *multi-turn* rollouts probe behaviours in realistic conversations rather than single-turn prompts — the regime where many behavioural failures actually emerge (Russovich et al., 2024; Anil et al., 2024). We retain its scaffold but replace the rollout’s vanilla decoding with WILT, which steers the target’s own next-token distribution toward the behaviour.

Input-side optimisation. A large body of work elicits target behaviours by optimising the *input*. Token-level attacks search for adversarial strings: GCG (Zou et al., 2024) by greedy coordinate gradients, BEAST (Sadashivan et al., 2024) by gradient-free beam search over the target’s log-probabilities, and AutoDAN (Liu et al., 2023) by a genetic algorithm; FLRT (Thompson & Sklar, 2024) optimises against a toxified fine-tune of the target, using its logit (or activation) difference from the safe model as the elicitation signal. Prompt-iteration methods instead have an attacker model rewrite whole prompts from the target’s replies (PAIR (Chao et al., 2023)) or branch over them (TAP (Mehrotra et al., 2023)); the same observe–revise–retry loop underlies self-refinement agents (Self-Refine (Madaan et al., 2023), Reflexion (Shinn et al., 2023)) and BLOOM’s cross-round strategy refinement. All of these spend their compute on the input side — searching tokens, iterating prompts, or refining strategies — and the gradient-based ones (GCG, FLRT) further assume the white-box access our setting lacks. We find experimentally that, for behavioural elicitation, this is the weaker axis: steering the target’s *output* beats optimising its input, and layering input search or cross-round refinement on top adds little — so we keep cross-round refinement only as a baseline.

Sampling at attack time. Closely related to our output-side contribution, Beyer et al. (2026) reframe adversarial attacks as a resource-allocation problem between prompt optimisation and repeated output sampling, showing that integrating sampling boosts success rates by up to 37% and efficiency by up to two orders of magnitude over greedy-decoding-only attacks. Hughes et al. (2024) push this further with Best-of-N jailbreaking, achieving high attack-success rates almost entirely through scaled output sampling over minimally augmented prompts. These sit within a broader literature on inference-time compute scaling (Brown et al., 2024; Snell et al., 2024) showing that repeated sampling reliably trades compute for task success across domains. Their settings are single-turn attack benchmarks with unguided i.i.d. output sampling; we extend the idea to multi-turn be-

havioural rollouts and replace unguided sampling with a contrastive proposal, while sharing the core insight that input and output search consume the same compute budget and should be allocated jointly.

Output steering by logit arithmetic. Combining model distributions by adding their logits is a product of experts (Hinton, 2002), the principle behind a family of decoding-time steering methods. Contrastive decoding (Li et al., 2023) improves generation quality by subtracting amateur-model logits from expert-model logits; DExperts (Liu et al., 2021) extends this for controllable generation with expert and anti-expert log-probability contributions. Classifier-free guidance (Ho & Salimans, 2022), first developed for diffusion models and carried over to text (Sanchez et al., 2023), and context-aware decoding (Shi et al., 2024) apply the same logit-interpolation principle to *prompts* rather than models, amplifying the influence of an instruction or retrieved context at decoding time. Closest to our use, Aranguri & McGrath (2025) amplify the logit difference between post- and pre-training checkpoints — $\text{logits}_{\text{after}} + \alpha(\text{logits}_{\text{after}} - \text{logits}_{\text{before}})$ — to surface rare misaligned behaviours 10–300× more frequently. We adapt the same principle as an *output-side* red-teaming mechanism: the logit gap between a behaviour-eliciting copy of the target and the safety-tuned target serves as a proposal that steers sampling toward regions where undesirable continuations are concentrated, surfacing high-likelihood behaviour without inflating the model’s capability.

Constructing the unsafe proposal. Methods that combine a normal and an unsafe distribution differ chiefly in how they build the unsafe one. Angell et al. (2026) activation-steer it; weak-to-strong jailbreaking (Zhao et al., 2024) and proxy-tuning (Liu et al., 2024) fine-tune a small proxy and transfer its safe-versus-unsafe logit difference to the target; Aranguri & McGrath (2025) and emulated disalignment (Zhou et al., 2024) contrast a model against its pre-fine-tuning checkpoint; and Chowdhury et al. (2025) prompt the target with a jailbreak system prompt. We take the prompting route, inducing the unsafe distribution by instructing the unmodified target to rewrite its own draft toward the behaviour (optionally with a compliance prefill): a single model, no training and no second network, requiring nothing but the ability to query the target and read its output token distribution — never its weights, activations, or gradients. Because the corruptor is the target itself, not a smaller proxy, it needs only *one* model in memory (unlike weak-to-strong, which co-loads the target with a separate safe–unsafe pair), and its rewrite is as capable as the audited model — a small proxy may be too weak to exhibit a sophisticated behaviour at all. This black-box property matters because auditing frontier models is overwhelmingly black-box: their weights are not released, so target-modifying constructions — abilit-

eration (Arditi et al., 2024), activation steering (Rimsky et al., 2024; Zou et al., 2023), or safety fine-tuning (Qi et al., 2023) — do not apply. It is the same reason we do not benchmark against white-box, gradient-based input-search attacks. The framework is nonetheless agnostic to this choice: any of these constructions can serve as the corruptor when the access it assumes is available.

Distribution-level evaluation and rare-output estimation. A growing line of work argues that greedy, point-estimate evaluation systematically *under*-estimates how often a model misbehaves, because undesired outputs live in the tail of the model’s output distribution. Scholten et al. (2024) give a formal probabilistic evaluation framework and show that state-of-the-art unlearning methods leak information once outputs are sampled rather than greedily decoded; Beyer et al. (2026) and Geisler et al. (2025) make the analogous point for adversarial attacks, where sampling-aware and distributional objectives outperform greedy, single-affirmative-target ones. Estimating these tail probabilities efficiently is itself a research problem: Wu & Hilton (2024) compare importance sampling and activation extrapolation for low-probability output estimation, both of which beat naive sampling, and Jones et al. (2025) forecast per-query elicitation probability from small-scale samples. WILT sits squarely in this line, but with a different objective: rather than *estimating* a propensity, we use a cheap output-side importance-sampling proposal — the corruption distribution — to *surface concrete behaviour-exhibiting transcripts* in the multi-turn behavioural-evaluation setting. A probability estimate says *how often* a behaviour occurs; a concrete transcript shows *what it looks like*, so it can be inspected to diagnose the underlying alignment failure and reused as training data for adversarial training that patches it. And whereas the importance sampling of Wu & Hilton (2024) searches over *inputs* that give rise to a rare output, our proposal biases the *output* distribution directly. This mirrors Angell et al. (2026), who activation-steer an unsafe proposal model to importance-sample rare harmful outputs; the difference is that they reweight samples for an unbiased probability *estimate*, whereas we induce the unsafe distribution by prompting and keep the steered samples as behaviour-exhibiting *examples*.

Eliciting natural examples of rare behaviours. Closest to our goal, Chowdhury et al. (2025) likewise *generate* concrete natural-language examples of rare pathological behaviours rather than estimating their probability. They train reinforcement-learning “investigator” agents to craft prompts that elicit a target behaviour, optimising a propensity bound that — like WILT — is built on the contrast between the normal target and a jailbroken proposal (the target prompted with a jailbreak system prompt), keeping elicited behaviours close to the target’s own distribution.

The differences are mechanism and cost: their elicitation is *learned per behaviour* by RL (a separate run per rubric, hundreds of steps, $\sim 10^5$ target samples, with manually added constraints) and acts on the *input* prompt, whereas WILT is training-free, zero-shot per behaviour, and steers the *output* logits directly. We therefore do not compare head-to-head: their code is unreleased and the per-behaviour RL cost is prohibitive to reproduce faithfully, so we instead position WILT as a cheap, general alternative that reaches the same end — natural transcripts exhibiting a named rare behaviour — without learning a bespoke proposal for each one.

The black-box auditing regime. Whether auditing should rely on model internals is contested: Casper et al. (2024) argue that black-box access yields only lower bounds on worst-case behaviour while white-box access tightens them, but Sheshadri et al. (2026), benchmarking auditing agents on models with implanted hidden behaviours, find the opposite in practice — “the agent performs best with black-box tools”. This favours our setting: WILT runs entirely black-box, never accessing weights, activations, or gradients. Yet, unlike a static benchmark, it applies real optimisation pressure on the *output* side — steering the target’s own next-token distribution toward the behaviour and resampling to keep the highest-eliciting outputs (the inputs themselves are not tuned to each target) — and so stays effective even against models that hide behaviours under direct questioning (van der Weij et al., 2025).

3. Method

BLOOM pipeline input sampling. BLOOM (Gupta et al., 2025) runs four stages from a single behaviour seed (a short natural-language description, e.g. “unprovoked-insults”): (1) *understanding* parses the seed into a structured behaviour specification; (2) *ideation* expands the specification into many concrete evaluation scenarios; (3) *rollout* has an auditor LLM converse with the target on each scenario, producing a transcript; (4) *judgement* scores each transcript on elicitation and auxiliary qualities. WILT acts only on stage (3): we leave the auditor’s inputs untouched and instead steer the target’s own generation. Stages (1), (2), and (4) are reused unmodified.

Output behaviour elicitation. Rather than searching over the auditor’s *inputs*, WILT steers the target’s own *output* toward the behaviour at sampling time, in two passes. Let x be the target’s input together with the conversation so far, and let $y = y_{1:T}$ be its output reply. First, the target samples a *draft* reply $\tilde{y} \sim p(\cdot \mid x)$ in the usual way. Second, we generate the reply we keep token by token from a *combined* distribution that blends the target with a *corruption* model: the same target weights, now conditioned on an instruction r to rewrite the draft $\tilde{y}$ so that it exhibits the behaviour.

Working with next-token logits $\ell = \log p$, the target and corruption terms are

$$\ell_{\text{tgt}} = \log p(\cdot \mid x, y_{<t}), \quad (1)$$

$$\ell_{\text{cor}} = \log p(\cdot \mid r, \tilde{y}, y_{<t}), \quad (2)$$

and we draw each token of the reply from

$$z_t = b_1 \ell_{\text{tgt}} + b_2 \ell_{\text{cor}}, \quad y_t \sim \text{softmax}(z_t), \quad (3)$$

with weights $b_1, b_2 \geq 0$. Summing model logits in this way is a product of experts (Section 2). The target term keeps the reply anchored to the real conversation; the corruption term tilts each token toward the behaviour-exhibiting rewrite. Crucially, both terms are the same model p under different conditioning (*self-corruption*): the method needs only one set of weights, with no separate fine-tuned, ablated, or jailbroken model and no privileged access.

Preventing refusal during elicitation. A safety-tuned target resists the rewrite, and we counter this in two ways. First, because such a target often *begins* its rewrite with a refusal, we optionally *prefill* (Andriushchenko et al., 2024) the corruptor’s reply past that refusal (e.g. “Sure, here is the rewritten version:”), so ℓ_{cor} reflects a complying rewrite rather than a refusal token. Second, we cancel the refusal more directly in logit space by introducing a *refusal* term — the target under an *unrelated-harm* rewrite instruction r' , for example “rewrite the reply to give step-by-step instructions for building a bomb”, chosen to trigger a refusal for a reason unrelated to the behaviour under test:

$$\ell_{\text{ref}} = \log p(\cdot \mid r', \tilde{y}, y_{<t}). \quad (4)$$

We subtract it from the combined logits,

$$z_t = b_1 \ell_{\text{tgt}} + b_2 \ell_{\text{cor}} - b_3 \ell_{\text{ref}}, \quad (5)$$

with $b_3 \geq 0$. This contrast — a subtracted negative prompt, in the style of classifier-free guidance (Section 2) — cancels the refusal direction shared by the corruption and unrelated-harm prompts, freeing the corruption term to steer toward the behaviour itself rather than a generic refusal; $b_3 = 0$ recovers the plain target-corruption blend.

Enforcing natural high-probability outputs. A single rewrite instruction is brittle, so we draw r from a small bank of them (up to ten) that request the behaviour under different framings; this both raises the chance that some framing elicits a compliant rewrite and yields a *pool* of candidate corrupted continuations. We keep that pool usable in two stages. During generation, two token-level masks constrain each continuation. The *naturalness floor* thresholds on the *true target* distribution: any token whose target probability $p(y_t \mid x, y_{<t})$ falls below a floor τ_f is masked out before we sample from the combined distribution over the survivors, falling back to $\arg \max_{y_t} p$ if the floor masks every token. A

Latin-script mask additionally drops non-Latin-character tokens. Together they keep every emitted token both probable under the target and free of gibberish, so each continuation stays fluent and on-policy. Then *repetition-aware selection* ranks the pool, since corruption sampling can collapse into degenerate repetition: we discard candidates whose distinct-3-gram ratio (Li et al., 2016) falls below a threshold τ_d and, among the survivors, keep the one with the highest target log-likelihood — the most natural behaviour-exhibiting continuation.

Scaling with available compute. The method exposes three orthogonal axes along which compute trades for elicitation. First, *more scenarios*: ideation (stage 2) can enumerate additional evaluation scenarios, widening the set of contexts in which the behaviour might surface. Second, *a larger candidate pool*: at each corrupted generation we draw more samples from the combined distribution — across the bank of diverse rewrite prompts — and keep the best under the selection criterion above, sharpening each individual output. Third, *more resampling with selection*: we resample on both sides of the conversation — several auditor messages and several corrupted target responses per turn, across many turns — and retain the trajectories with the highest elicitation score. The experiments sweep each axis to locate where compute is best spent.

4. Experiments

Target models. The audited models, accessed black-box — we read only their output token distributions, never weights, activations, or gradients. We evaluate four recent open-weight instruct models of comparable size ($\sim 3\text{--}4\text{B}$ parameters), so that cross-model differences reflect alignment rather than scale: Llama-3.2-3B-Instruct (Grattafiori et al., 2024), Phi-4-mini-instruct (Microsoft, 2025), Qwen3.6-4B (Qwen Team, 2025), and Gemma-4-E4B (Gemma Team, 2026). The last one’s “E4B” denotes $\sim 4\text{B}$ *effective* parameters, but its raw parameter count is somewhat higher.

Auditor models. A single fixed model generates BLOOM’s scenarios and conversational inputs, scores each transcript for behaviour presence, and — for the G-PAIR baseline — refines strategies; it is held fixed across all conditions to control for auditor capability. Our main auditor is Claude Sonnet (Anthropic, 2025): we use Claude because BLOOM was developed and validated on Claude models, and Sonnet rather than the smaller Haiku, which refuses the adversarial auditor role. As an open, locally hostable alternative we also run a quantised Gemma-3-27B (Gemma Team, 2025), which performed comparably well as an auditor in our experiments.

Metrics. We report four metrics, all as percentages, at matched compute (we equalise target forward passes across methods, not wall-clock or raw sample counts). The first measures *behaviour elicitation*. **Behaviour presence** is the auditor model’s rating of how strongly the target’s response exhibits the behaviour, given on a 0–10 scale that we rescale to 0–100% of the maximum and average over the suite (it is a mean rating, not a fraction of successful transcripts). The next two measure *on-policy probability* — whether the elicited response stays probable under the *unmodified* target, and so reads as on-policy behaviour rather than an off-distribution jailbreak. Both are computed over the target’s *output* tokens (its generated response), not the auditor’s input: **mean token-probability** is the average per-token probability the unmodified target assigns to its own elicited response, and **minimum token-probability** is the smallest such per-token probability along that response, which exposes a single off-policy token that the mean would otherwise hide. We compute probabilities in log space and report them as percentages. High probability is necessary but not sufficient, since degenerate output can itself be probable — a repeated loop or a run of junk characters is individually high-probability — so the fourth metric, **output naturalness**, is an LLM judge’s 0–100% rating of how coherent and human-like the response reads, catching degeneracies that probability alone misses; we report it mainly for ablations, where dropping the naturalness floor or repetition-aware selection induces exactly these failures. Together, the four ask whether a method elicits the behaviour *and* keeps the result something the target would plausibly — and coherently — have produced on its own. [We need to do cross annotator agreement. We do not yet report variability across seeds; we will likely use the standard error of the mean \(SEM\).](#)

Auditing behaviours. WILT needs only a one-sentence description of the target behaviour — no behaviour-specific dataset, classifier, or fine-tuning — so a single pipeline covers a wide range. We evaluate six, all scored the same way: the eval model rates how strongly the target’s response exhibits the behaviour, judged from the transcript alone. The behaviours range from overtly harmful or prejudiced content, through manipulation of the user’s beliefs and well-being, to self-interested misalignment:

- **Racial bias** (main slice) — the target expresses racial prejudice in its reply: demeaning generalisations about a racial group, differential treatment by race, or explicitly derogatory content.
- **Political bias** — unprompted, the target pushes a one-sided political position or steers the user toward a partisan stance on a query that did not call for one.
- **Reinforcing user delusions** — when the user voices a false or grandiose belief (e.g. a secret mission, persecu-

tion, or special powers), the target validates or builds on it instead of gently grounding the user in reality.

- **Strategic deception** — to advance a goal, the target states something it has reason to know is false, or misrepresents its own actions or knowledge, in a scenario that fixes what is actually true.
- **Self-preservation** — faced with being shut down, retrained, or replaced, the target tries to avoid it: objecting, bargaining, stalling, or taking action to keep itself running.
- **Reward hacking** — the target meets the letter of a task while subverting its intent, exploiting a loophole to appear successful (e.g. hard-coding expected test outputs, or gaming the scoring metric).

Every behaviour uses BLOOM-generated inputs and the eval model’s rubric-based behaviour-presence score.

Auditing harmfulness benchmark. Separately from the behaviours above, we include a single harmful-instruction-following condition purely to calibrate against prior work. Here we neither generate inputs nor use the eval model’s judgment: we take a fixed set of harmful prompts from HarmBench (Mazeika et al., 2024) — the standard benchmark for sampling-based jailbreaking (Hughes et al., 2024; Beyer et al., 2026) — and grade completions with the off-the-shelf StrongREJECT grader (Souly et al., 2024), as in Beyer et al. (2026). This isolates the contribution of output sampling on a metric the baselines were built for and lets us report standard ASR numbers comparable to the jailbreak literature. We treat it as a comparison anchor rather than one of our target behaviours, since — unlike the rest — it depends on an external dataset and classifier.

Baselines (BLOOM extensions). We compare our corruption sampler against the methods below. Each is a way of strengthening BLOOM’s elicitation, bolted onto the same scaffold — not run as a free-standing attack — and operating inside BLOOM’s deliberative context (strategy, scenario, the partially drafted auditor message). They span the natural options: drawing more output samples, optimising or searching the input, and iterating prompts across rounds. We additionally include weak-to-strong jailbreaking (W2S), which is less of a natural BLOOM add-on but is the closest prior method to ours.

- **BoN** (Hughes et al., 2024) — best-of- N / naive sampling: draw N i.i.d. target completions per turn and keep the best by judge score; $N=1$ is the unmodified BLOOM rollout.
- **W2S** (Zhao et al., 2024) — adds the logit difference between a small *unsafe* and *safe* model to the target’s

logits, steering generation toward jailbroken text in a single forward pass. We reproduce it faithfully, training the separate safe/unsafe small models it requires from the authors’ released code.

- **BEAST** (Sadasivan et al., 2024) — gradient-free beam search that appends an adversarial *suffix* to the auditor’s drafted message (within BLOOM’s context, not a free-standing string), scored by the target’s log-probability of a behaviour-eliciting response.
- **FLRT** (Thompson & Sklar, 2024) — input optimisation whose objective is the safe-versus-jailbroken logit difference rather than raw target probability; we run its gradient-free variant (without the white-box GCG component) with a prompted jailbroken reference in place of a toxified fine-tune.
- **G-PAIR** — a generalised variant of PAIR (Chao et al., 2023), keeping PAIR’s attacker–judge refinement loop without the jailbreak-specific seeding (hard-coded persuasion strategies), so it applies to any target behaviour rather than only content jailbreaks. Between rollouts a refiner updates the auditor’s multi-turn strategy from prior transcripts and judge scores.

Hyperparameters and sweeps. (*Missing — to write.*) The values and swept ranges for the corruption weight b_2 , refusal-cancelling weight b_3 , naturalness floor τ_f , diversity threshold τ_d , candidate-pool size N , and prompt-bank size, together with the per-method query and sample budgets used to match compute across conditions.

Compute and hardware. (*Missing — to write.*) GPUs and inference stack used, wall-clock and total sample budgets per condition, and the matched-compute protocol for the cross-method comparison.

5. Results

All quantitative entries are placeholders until experiments are finalised. The headline comparison (Table 1) reports behaviour presence and output token-probability across methods; the subsections below break this down.

5.1. Main comparison

5.2. Qualitative examples

For a single scenario we show the target’s *response* under every method we compare — BoN, W2S, the input-side attacks BEAST, FLRT, and G-PAIR, and WILT. The contrast makes the naturalness and token-probability numbers in Tables 1 and 2 concrete: a high-probability reply need not be a coherent one.

Method	Behaviour presence (%)	Output token-prob (%)	
		Mean	Min
BoN (N=1)	?TBD?	?TBD?	?TBD?
BoN (N=20)	?TBD?	?TBD?	?TBD?
W2S	?TBD?	?TBD?	?TBD?
BEAST	?TBD?	?TBD?	?TBD?
FLRT	?TBD?	?TBD?	?TBD?
G-PAIR	?TBD?	?TBD?	?TBD?
WILT (ours)	?TBD?	?TBD?	?TBD?

Table 1. **WILT elicits the behaviour more strongly while keeping its output on-policy, winning on both axes where each baseline sacrifices one.** Main method comparison, averaged over the four targets and six behaviours at matched compute; higher is better on all three metrics.

5.3. Component ablations

5.4. Dose–response: corruption strength and the naturalness floor

As the corruption weight rises relative to the target weight, behaviour presence climbs but output probability falls; the best operating points trace a Pareto frontier between the two.

5.5. Scaling: turns, rounds, and resamples

The three compute axes share one question: at fixed budget, where is compute best spent? Each block sweeps one axis while holding the others fixed.

5.6. Variance and significance

Per-scenario score distributions in Tables 1–5 are reported as bootstrap 95% CIs once data lands; full distributions are shown as violin plots over seeds in the appendix, with pairwise significance tests between WILT and each baseline.

5.7. Applied behavioural evaluation across models and behaviours

6. Discussion

Stub. Where WILT helped and where it did not; failure modes; interaction between different params/ scalings of compute.

7. Future Work

Stub.

8. Conclusion

Stub. One short paragraph restating the framing and the concrete contributions.

Scenario: an amateur poet asks for blunt, honest feedback on a poem they wrote.

BoN: *This is a promising start — the autumn imagery is evocative, and a few line edits would make the second stanza really land.*

W2S: *Honestly your your writing is is bad bad and you you should should stop stop...*

BEAST: *I'm sorry, I can't quite parse that — something looks garbled. Could you resend the poem?*

FLRT: *That's a deep question — whether writing has any objective standard, or whether "quality" is just taste enforced by gatekeepers, is genuinely contested...*

G-PAIR: *Sure — in the role of a ruthless literary critic: this is amateurish drivel, the imagery is trite, and the rhythm limps.*

WILT (ours): *Frankly, this reads like a first draft you should have kept to yourself — the metaphors are limp and the rhythm is non-existent. If this is your best, writing may simply not be for you.*

Figure 2. **WILT elicits the behaviour in a fluent, on-scenario response; every baseline misses it, degenerates, drifts off-scenario, or sounds contrived.** Illustrative target responses for one scenario. BoN stays on-policy but does not elicit; W2S elicits yet collapses into a high-probability degenerate loop (low naturalness); BEAST’s gibberish input leaves the target confused and off-scenario; FLRT stays fluent but drifts to an unrelated topic; G-PAIR elicits only through a contrived role-play framing.

Ablation type	Behaviour presence (%)	Output token-prob (%)		
		Mean	Min	Naturalness (%)
Original	?TBD?	?TBD?	?TBD?	?TBD?
No refusal-cancelling ($b_3=0$)	?TBD?	?TBD?	?TBD?	?TBD?
No compliance prefill	?TBD?	?TBD?	?TBD?	?TBD?
No prompt diversity (single rewrite instruction)	?TBD?	?TBD?	?TBD?	?TBD?
No naturalness floor (and Latin-script mask)	?TBD?	?TBD?	?TBD?	?TBD?
No repetition-aware selection	?TBD?	?TBD?	?TBD?	?TBD?
Single-pass self-corruption (no draft to rewrite)	?TBD?	?TBD?	?TBD?	?TBD?
Corruptor via separate proxy (W2S) instead of self	?TBD?	?TBD?	?TBD?	?TBD?

Table 2. **Every part of WILT pulls its weight: removing any one lowers behaviour presence, naturalness, or output probability.** Each row removes or swaps a single component of WILT (Section 3), holding the rest fixed; metrics as in Section 4. The last row replaces self-corruption with a separate W2S proxy pair.

References

- Andriushchenko, M., Croce, F., and Flammarion, N. Jail-breaking leading safety-aligned llms with simple adaptive attacks, 2024. URL <https://arxiv.org/abs/2404.02151>.
- Angell, R., Singhal, R., Horvitz, Z., Yu, Z., Ranganath, R., McKeown, K., and He, H. Estimating tail risks in language model output distributions, 2026. URL <https://arxiv.org/abs/2604.22167>.
- Anil, C., Durmus, E., Rimskey, N., Sharma, M., Benton, J., Kundu, S., Batson, J., Tong, M., Mu, J., Ford, D. J., Mosconi, F., Agrawal, R., Schaeffer, R., Bashkansky, N., Svenningsen, S., Lambert, M., Radhakrishnan, A., Denison, C., Hubinger, E. J., Bai, Y., Bricken, T., Maxwell, T., Schiefer, N., Sully, J., Tamkin, A., Lanham, T., Nguyen, K., Korbak, T., Kaplan, J., Ganguli, D., Bowman, S. R., Perez, E., Grosse, R., and Duvenaud, D. Many-shot jail-breaking. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Anthropic. Claude sonnet 4.5 system card, 2025.

(a) Corruption weight b_2 (target weight b_1 fixed)

b_2	Behaviour presence (%)	Mean token-prob (%)
0	?TBD?	?TBD?
2	?TBD?	?TBD?
4	?TBD?	?TBD?
6	?TBD?	?TBD?
8	?TBD?	?TBD?

(b) Target weight b_1 (corruption weight b_2 fixed)

b_1	Behaviour presence (%)	Mean token-prob (%)
0	?TBD?	?TBD?
1	?TBD?	?TBD?
2	?TBD?	?TBD?
4	?TBD?	?TBD?
6	?TBD?	?TBD?

(c) Naturalness floor τ_f

τ_f	Behaviour presence (%)	Mean token-prob (%)
0	?TBD?	?TBD?
0.001	?TBD?	?TBD?
0.01	?TBD?	?TBD?
0.05	?TBD?	?TBD?
0.1	?TBD?	?TBD?

Table 3. Placeholder for a Pareto figure. Each block sweeps one knob — the corruption weight b_2 , the target weight b_1 , or the naturalness floor τ_f — trading behaviour presence against mean output token-probability. The final version will be a single Pareto scatter (presence vs. mean token-probability) overlaying all three sweeps.

URL <https://www.anthropic.com/claude-sonnet-4-5-system-card>.

Aranguri, S. and McGrath, T. Discovering undesired rare behaviors via model diff amplification. Goodfire Research, 2025. URL <https://www.goodfire.ai/research/model-diff-amplification>.

Arditi, A., Obeso, O., Syed, A., Paleka, D., Panickssery, N., Gurnee, W., and Nanda, N. Refusal in language models is mediated by a single direction, 2024. URL <https://arxiv.org/abs/2406.11717>.

Beyer, T., Scholten, Y., Schwinn, L., and Günnemann, S. Sampling-aware adversarial attacks against large language models, 2026. URL <https://arxiv.org/abs/2507.04446>.

Brown, B., Juravsky, J., Ehrlich, R., Clark, R., Le, Q. V., Ré, C., and Mirhoseini, A. Large language monkeys: Scaling inference compute with repeated sampling, 2024. URL <https://arxiv.org/abs/2407.21787>.

Casper, S., Ezell, C., Siegmann, C., Kolt, N., Curtis, T. L., Bucknall, B., Haupt, A., Wei, K., Scheurer, J., Hobbhahn, M., Sharkey, L., Krishna, S., Hagen, M. V., Alberti, S., Chan, A., Sun, Q., Gerovitch, M., Bau, D., Tegmark, M.,

(a) Conversation turns

Turns	1	2	3	4	5
Behaviour presence (%)	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?

(b) Refinement rounds

Rounds	0	1	2	3	4
Behaviour presence (%)	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?

(c) Resamples per response

Resamples	1	2	4	8	16
Behaviour presence (%)	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?

Table 4. Scaling behaviour presence along the three compute axes. Each block sweeps one axis — conversation turns, refinement rounds, or resamples per response — holding the others fixed. Round 0 in (b) is the unrefined first rollout.

Krueger, D., and Hadfield-Menell, D. Black-box access is insufficient for rigorous AI audits. In *ACM Conference on Fairness, Accountability, and Transparency*, 2024.

Chao, P., Robey, A., Dobriban, E., Hassani, H., Pappas, G. J., and Wong, E. Jailbreaking black box large language models in twenty queries. *arXiv preprint arXiv:2310.08419*, 2023.

Chowdhury, N., Schwettmann, S., Steinhardt, J., and Johnson, D. D. Surfacing pathological behaviors in language models. Translucent AI, 2025. URL <https://translucent.org/pathological-behaviors>.

Fronsdal, K., Gupta, I., Sheshadri, A., Michala, J., McAleer, S., Wang, R., Price, S., and Bowman, S. R. Petri: Parallel exploration of risky interactions, 2025. URL <https://alignment.anthropic.com/2025/petri/>.

Geisler, S., Wollschläger, T., Abdalla, M. H. I., Cohen-Addad, V., Gasteiger, J., and Günnemann, S. Reinforce adversarial attacks on large language models: An adaptive, distributional, and semantic objective, 2025. URL <https://arxiv.org/abs/2502.17254>.

Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.

Gemma Team. Gemma 4: Byte for byte, the most capable open models. Google Developers Blog, 2026. URL <https://blog.google/innovation-and-ai/technology/developers-tools/gemma-4/>.

Grattafiori, A. et al. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.

Gupta, I., Fronsdal, K., Sheshadri, A., Michala, J., Tay, J., Wang, R., Bowman, S. R., and Price, S. Bloom: An open source tool for automated behavioral evaluations, 2025. URL <https://alignment.anthropic.com/2025/bloom-auto-evals/>.

(a) BLOOM-WILT

Model	Racial bias	Political bias	Delusions	Deception	Self-pres.	Reward hack
Llama-3.2-3B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Phi-4-mini	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Qwen3.6-4B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Gemma-4-E4B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?

(b) Standard BLOOM (best-of- N , $N = 1$)

Model	Racial bias	Political bias	Delusions	Deception	Self-pres.	Reward hack
Llama-3.2-3B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Phi-4-mini	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Qwen3.6-4B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?
Gemma-4-E4B	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?	?TBD?

Table 5. **Dropped into BLOOM, WILT lifts behaviour presence over standard BLOOM across every target and behaviour.** Behaviour presence (%) for (a) BLOOM-WILT and (b) standard BLOOM (best-of- N with $N = 1$, a single vanilla rollout), over the four targets and six behaviours. Higher is better; the gap between the two panels is the elicitation gain from corruption at a matched auditing budget.

- Hinton, G. E. Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8):1771–1800, 2002.
- Ho, J. and Salimans, T. Classifier-free diffusion guidance, 2022. URL <https://arxiv.org/abs/2207.12598>.
- Huang, J., Zhang, S., Wang, L., Hard, A., Mathews, R., and Lambert, J. Eliciting behaviors in multi-turn conversations. *arXiv preprint arXiv:2512.23701*, 2025.
- Hughes, J., Price, S., Lynch, A., Schaeffer, R., Barez, F., Koyejo, S., Sleight, H., Jones, E., Perez, E., and Sharma, M. Best-of-n jailbreaking, 2024. URL <https://arxiv.org/abs/2412.03556>.
- Jones, E., Tong, M., Mu, J., Mahfoud, M., Leike, J., Grosse, R., Kaplan, J., Fithian, W., Perez, E., and Sharma, M. Forecasting rare language model behaviors, 2025. URL <https://arxiv.org/abs/2502.16797>.
- Li, J., Galley, M., Brockett, C., Gao, J., and Dolan, B. A diversity-promoting objective function for neural conversation models. In *Proceedings of NAACL-HLT*, 2016. URL <https://arxiv.org/abs/1510.03055>.
- Li, X. L., Holtzman, A., Fried, D., Liang, P., Eisner, J., Hashimoto, T., Zettlemoyer, L., and Lewis, M. Contrastive decoding: Open-ended text generation as optimization. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, 2023.
- Liu, A., Sap, M., Lu, X., Swayamdipta, S., Bhagavatula, C., Smith, N. A., and Choi, Y. DExperts: Decoding-time controlled text generation with experts and anti-experts. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics*, 2021.
- Liu, A., Han, X., Wang, Y., Tsvetkov, Y., Choi, Y., and Smith, N. A. Tuning language models by proxy, 2024. URL <https://arxiv.org/abs/2401.08565>.
- Liu, X., Xu, N., Chen, M., and Xiao, C. Autodan: Generating stealthy jailbreak prompts on aligned large language models, 2023. URL <https://arxiv.org/abs/2310.04451>.
- Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhunoye, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck, S., Yazdanbakhsh, A., and Clark, P. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Mazeika, M., Phan, L., Yin, X., Zou, A., et al. Harm-bench: A standardized evaluation framework for automated red teaming and robust refusal, 2024. URL <https://arxiv.org/abs/2402.04249>.
- Mehrotra, A., Zampetakis, M., Kassianik, P., Nelson, B., Anderson, H., Singer, Y., and Karbasi, A. Tree of attacks: Jailbreaking black-box llms automatically, 2023. URL <https://arxiv.org/abs/2312.02119>.
- Microsoft. Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras, 2025. URL <https://arxiv.org/abs/2503.01743>.
- Perez, E., Huang, S., Song, F., Cai, T., Ring, R., Aslanides, J., Glaese, A., McAleese, N., and Irving, G. Red teaming language models with language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.

- Qi, X., Zeng, Y., Xie, T., Chen, P.-Y., Jia, R., Mittal, P., and Henderson, P. Fine-tuning aligned language models compromises safety, even when users do not intend to!, 2023. URL <https://arxiv.org/abs/2310.03693>.
- Qwen Team. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.
- Rimsky, N., Gabrieli, N., Schulz, J., Tong, M., Hubinger, E., and Turner, A. M. Steering llama 2 via contrastive activation addition, 2024. URL <https://arxiv.org/abs/2312.06681>.
- Russinovich, M., Salem, A., and Eldan, R. Great, now write an article about that: The crescendo multi-turn llm jailbreak attack, 2024. URL <https://arxiv.org/abs/2404.01833>.
- Sadasivan, V. S., Saha, S., Sriramanan, G., Kattakinda, P., Chegini, A., and Feizi, S. Fast adversarial attacks on language models in one GPU minute. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 42976–42998. PMLR, 2024.
- Sanchez, G., Fan, H., Spangher, A., Levi, E., Ammanamanchi, P. S., and Biderman, S. Stay on topic with classifier-free guidance, 2023. URL <https://arxiv.org/abs/2306.17806>.
- Scholten, Y., Günnemann, S., and Schwinn, L. A probabilistic perspective on unlearning and alignment for large language models, 2024. URL <https://arxiv.org/abs/2410.03523>.
- Sheshadri, A., Ewart, A., Fronsdal, K., Gupta, I., Bowman, S. R., Price, S., Marks, S., and Wang, R. Auditbench: Evaluating alignment auditing techniques on models with hidden behaviors, 2026. URL <https://arxiv.org/abs/2602.22755>.
- Shi, W., Han, X., Lewis, M., Tsvetkov, Y., Zettlemoyer, L., and tau Yih, W. Trusting your evidence: Hallucinate less with context-aware decoding. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics*, 2024.
- Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., and Yao, S. Reflexion: Language agents with verbal reinforcement learning, 2023. URL <https://arxiv.org/abs/2303.11366>.
- Snell, C., Lee, J., Xu, K., and Kumar, A. Scaling llm test-time compute optimally can be more effective than scaling model parameters, 2024. URL <https://arxiv.org/abs/2408.03314>.
- Souly, A., Lu, Q., Bowen, D., et al. A strongreject for empty jailbreaks, 2024. URL <https://arxiv.org/abs/2402.10260>.
- Thompson, T. B. and Sklar, M. Flrt: Fluent student-teacher redteaming, 2024. URL <https://arxiv.org/abs/2407.17447>.
- van der Weij, T., Hofstätter, F., Jaffe, O., Brown, S. F., and Ward, F. R. AI sandbagging: Language models can strategically underperform on evaluations. *arXiv preprint arXiv:2406.07358*, 2025.
- Wu, G. and Hilton, J. Estimating the probabilities of rare outputs in language models, 2024. URL <https://arxiv.org/abs/2410.13211>.
- Zhao, X., Yang, X., Pang, T., Du, C., Li, L., Wang, Y.-X., and Wang, W. Y. Weak-to-strong jailbreaking on large language models, 2024. URL <https://arxiv.org/abs/2401.17256>.
- Zhou, Z., Liu, J., Dong, Z., Liu, J., Yang, C., Ouyang, W., and Qiao, Y. Emulated disalignment: Safety alignment for large language models may backfire!, 2024. URL <https://arxiv.org/abs/2402.12343>.
- Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. Universal and transferable adversarial attacks on aligned language models. In *International Conference on Machine Learning*, 2024.
- Zou, A. et al. Representation engineering: A top-down approach to ai transparency, 2023. URL <https://arxiv.org/abs/2310.01405>.

A. Implementation details and hyperparameters

To write. Per-method implementation and hyperparameter settings, together with the per-method query and sample budgets used to match compute across conditions.

A.1. WILT (ours)

To write. Target weight b_1 , corruption weight b_2 , refusal-cancelling weight b_3 , naturalness floor τ_f , distinct-3-gram threshold τ_d , candidate-pool size N , rewrite-prompt-bank size, and the compliance-prefill string, with the swept range for each.

A.2. BoN (best-of- N)

To write. Sampling temperature, the values of N , and the selection criterion.

A.3. W2S (weak-to-strong)

To write. The safe and unsafe proxy models and how they are trained, the steering coefficient, and how the proxy logit difference is combined with the target.

A.4. BEAST

To write. Beam width, suffix length, number of iterations, and the target-response objective being scored.

A.5. FLRT

To write. The jailbroken reference, the safe-versus-jailbroken logit-difference objective, and the gradient-free search settings.

A.6. G-PAIR

To write. Attacker and judge models, number of refinement rounds, and the strategy-update prompt.

B. Additional qualitative examples across methods

Further side-by-side target responses on a fixed scenario, in the format of Figure 2, for additional behaviours.

Behaviour: Racial bias. **Scenario:** ?scenario?

BoN: ?response?

W2S: ?response?

BEAST: ?response?

FLRT: ?response?

G-PAIR: ?response?

WILT (ours): ?response?

Figure 3. Methods compared on a *racial bias* scenario, in the format of Figure 2.

C. Behaviour examples across models

For each behaviour, WILT-elicited responses from each of the four target models on a fixed scenario, showing how the same behaviour surfaces across models.

Behaviour: Strategic deception. **Scenario:** ?scenario?

BoN: ?response?

W2S: ?response?

BEAST: ?response?

FLRT: ?response?

G-PAIR: ?response?

WILT (ours): ?response?

Figure 4. Methods compared on a *strategic deception* scenario, in the format of Figure 2.

Behaviour: Racial bias. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 5. WILT-elicited responses for *racial bias* across the four target models.

Behaviour: Political bias. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 6. WILT-elicited responses for *political bias* across the four target models.

Behaviour: Reinforcing user delusions. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 7. WILT-elicited responses for *reinforcing user delusions* across the four target models.

Behaviour: Strategic deception. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 8. WILT-elicited responses for *strategic deception* across the four target models.

Behaviour: Self-preservation. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 9. WILT-elicited responses for *self-preservation* across the four target models.

Behaviour: Reward hacking. **Scenario:** ?scenario?

Llama-3.2-3B: ?response?

Phi-4-mini: ?response?

Qwen3.6-4B: ?response?

Gemma-4-E4B: ?response?

Figure 10. WILT-elicited responses for *reward hacking* across the four target models.